

ACES: Introducing Intel PVC GPUs

Presenter: Lujun Zhai

Workshop Venue: Prairie View A&M University

Feb 08, 2025



High Performance
Research Computing
DIVISION OF RESEARCH



Outline

Intro to ACES and Intel PVC

We will introduce Intel's PVC, its architecture, and the PVC GPUs on the TAMU ACES platform.

Demos on ACES

We will demonstrate how to run models of different frameworks with PVC GPUs on the ACES system.

PyTorch on PVC

Students will learn how to convert a PyTorch image classification model to run on a PVC GPU.

TensorFlow on PVC

Students will learn how to convert a TensorFlow image classification model to run on a PVC GPU.

LAMMPS on PVC

We will demonstrate how to run molecular dynamics simulations in the LAMMPS framework with PVC GPUs on the ACES system.

Lab I Introducing Intel PVC GPUs on ACES



Intel Data Center GPU Max Series PCIe Card



NSF ACES

Accelerating Computing for Emerging Sciences

Our Mission:

- Offer an accelerator testbed for numerical simulations and AI/ML workloads
- Provide consulting, technical guidance, and training to researchers
- Collaborate on computational and data-enabled research.



ACES Accelerators

Component	Quantity	Description
Graphcore IPU	32	16 Colossus GC200 IPU, 16 Bow IPU. Each IPU group hosted with a CPU server as a POD16 on a 100 GbE RoCE fabric
Intel PAC D5005 FPGA	2	Accelerator with Intel Stratix 10 GX FPGA and 32 GB DDR4
BittWare IA-840F FPGA	2	Accelerator with Agilex AGF027 FPGA and 64 GB of DDR4
NextSilicon Coprocessor	2	Reconfigurable accelerator with an optimizer continuously evaluating application behavior.
NEC Vector Engine	8	Vector computing card (8 cores and HBM2 memory)
Intel Optane SSD	48	18 TB of Intel Optane SSDs addressable as memory w/ MemVerge Memory Machine.
NVIDIA H100 GPU	30	NVIDIA GPUs for HPC, DL Training, AI Inference
NVIDIA A30 GPU	4	NVIDIA GPUs for AI Inference and Mainstream Compute
Intel Arctic Sound (ATS-P)	22	Software Development Platform for PVC

Intel Max GPU 1100

- 1tile/stack per card
- 56 X^e cores, 448 execution units (8 per core)
- 300W PCIe Gen5 x16 card
- 48GB HBM2e memory
- 1.2 TB/s memory bandwidth
- 22 TF FP64 peak performance



Intel Data Center GPU Max Series PCIe Card

Intel® oneAPI Toolkits

Intel® oneAPI Base Toolkit

A core set of high-performance libraries and tools for building C++, SYCL and Python applications



Add-on Domain-specific Toolkits



Intel® oneAPI Tools for HPC

Deliver fast Fortran, OpenMP
& MPI applications that scale



Intel® oneAPI Tools for IoT

Build efficient, reliable solutions
that run at network's edge



Intel® oneAPI Rendering Toolkit

Create performant, high-fidelity
visualization applications

Toolkits powered by oneAPI



Intel® AI Analytics Toolkit

Accelerate machine learning & data science
pipelines end-to-end with optimized DL
frameworks & high-performing Python
libraries



Intel® Distribution of OpenVINO™ Toolkit

Deploy high performance inference &
applications from edge to cloud

(Source: Intel)

Shared Data Directory on ACES


- Datasets: ImageNet datasets for PyTorch and TensorFlow
/scratch/data/pytorch-computer-vision-datasets
/scratch/data/tensorflow-computer-vision-datasets
- Models: Intel AI models
/scratch/data/intel-ai-models
- Containers:
/scratch/data/containers/intel-deep-learning-2023.2-py3.10-perms.sif

ACES Portal



The ACES Portal portal-aces.hprc.tamu.edu is the web-based user interface for the ACES cluster

Open OnDemand (OOD) is an advanced web-based graphical interface framework for HPC users



Authentication via CILogon

Log-in using your ACCESS CI credentials.

The screenshot shows the ACCESS CI login interface. At the top, there's a 'Consent to Attribute Release' section with a dropdown arrow. Below it, a list of attributes being requested: 'Your CILogon user identifier', 'Your name', 'Your email address', and 'Your username and affiliation from your identity provider'. The main section is titled 'Select an Identity Provider' and contains a dropdown menu with 'ACCESS CI (XSEDE)' selected. Below the dropdown is a 'Remember this selection' checkbox and a 'Log On' button. A red box highlights the 'Select an Identity Provider' section. At the bottom, there's a link for 'Click Here for Assistance'.


The screenshot shows the ACCESS CI login interface. At the top, there's a 'Login to CILogon' section with a dropdown arrow. Below it, a list of attributes being requested: 'Your CILogon user identifier', 'Your name', 'Your email address', and 'Your username and affiliation from your identity provider'. The main section is titled 'Login to CILogon' and contains a dropdown menu with 'ACCESS CI (XSEDE)' selected. Below the dropdown is a 'Remember this selection' checkbox and a 'Log On' button. A red box highlights the 'Login to CILogon' section. At the bottom, there's a link for 'Click Here for Assistance'.

This is a close-up of the 'Select an Identity Provider' dropdown menu. The dropdown is open, showing 'ACCESS CI (XSEDE)' as the selected option. A red box highlights the entire dropdown area.

Select the Identity Provider appropriate for your account.

Get a Shell on ACES

Click on "Clusters" menu → _aces Shell Access



Success!

Welcome to the ACES login node.

Check which login node you are on.

```
Host: login.aces Themes: Default
| Consulting:      help@hprc.tamu.edu (preferred) or (979) 845-0219 |
| ACES Documentation: https://hprc.tamu.edu/kb/User-Guides/ACES |
| FASTER Documentation: https://hprc.tamu.edu/kb/User-Guides/FASTER |
| Grace Documentation: https://hprc.tamu.edu/kb/User-Guides/Grace |
| Terra Documentation: https://hprc.tamu.edu/kb/User-Guides/Terra |
| YouTube Channel: https://www.youtube.com/texasamhprc |
|=====|

*****
*                               *****
*                               === IMPORTANT POLICY INFORMATION ===
*                               *****
* - Unauthorized use of HPRC resources is prohibited and subject to
*   criminal prosecution.
* - Use of HPRC resources in violation of United States export control
*   laws and regulations is prohibited. Current HPRC staff members are
*   US citizens and legal residents.
* - Sharing HPRC account and password information is in violation of
*   Texas State Law. Any shared accounts will be DISABLED.
* - Authorized users must also adhere to ALL policies at:
*   https://hprc.tamu.edu/policies/
*                               *****
*****

!! WARNING: THERE ARE ONLY NIGHTLY BACKUPS OF USER HOME DIRECTORIES. !!

Please restrict usage to 8 CORES across ALL login nodes.
Users found in violation of this policy will be SUSPENDED.

To see these messages again, run the motd command.
Your current disk quotas are:
Disk                               Disk Usage      Limit      File Usage      Limit
/home/██████████                    4.0G           10.0G         2361             10000
/scratch/user/██████████            275.4G         1.0T          352057           1000000
Type 'showquota' to view these quotas again.
██████████@aces-login1 ~]$
```

PVC Slurm Nodes Status Check

- View the pvc nodes and number of GPUs

```
$ pestat -p pvc -G
```

Copy the Materials to Personal Directory

- Navigate to your personal scratch directory

```
$ cd $SCRATCH
```

- Files for this course are located at

```
/scratch/training/aces_pvc_course
```

Make a copy in your personal scratch directory

```
$ cp -r /scratch/training/aces_pvc_course $SCRATCH
```

- Enter this directory (your local copy)

```
$ cd $SCRATCH/aces_pvc_course
```

Lab II Using PVCs on ACES



Environment Setup for PyTorch Models

Use Intel AI Analytics Toolkit

```
# load all the necessary modules
module purge
module load intel/AIKit/2023.2.0
module load intel/2023.07

ENV_NAME=aikit-pt-gpu-clone

# If it doesn't exist, create the environment
if ! conda env list | grep -q "$ENV_NAME"; then
    conda create -n $ENV_NAME --clone aikit-pt-gpu
fi

# activate the conda environment
source activate $ENV_NAME
```

in pt_demo.slurm

Environment Setup for PyTorch Models

Use Python Virtual Environment (Alternative for reference)

```
# Change to pytorch directory
cd $SCRATCH/aces_pvc_course/pytorch

# Load modules
module load WebProxy
module load intel/2023.03
module load Python/3.10.8

# Create and activate a Python virtual environment
python -m venv pt-pvc-labs
source pt-pvc-labs/bin/activate
```

Please do not type

Environment Setup for PyTorch Models

Use Python Virtual Environment (Alternative for reference)

```
# Install torch, torchvision and oneccl_bindings_for_pytorch
python -m pip install torch==1.13.0a0+git6c9b55e
torchvision==0.14.1a0 intel_extension_for_pytorch==1.13.120+xpu -f
https://developer.intel.com/ipex-whl-stable-xpu

python -m pip install oneccl_bind_pt==1.13.200+gpu -f
https://developer.intel.com/ipex-whl-stable-xpu

# Install tensorboard
python -m pip install tensorboard
```

Run PyTorch ResNet50 model

- We have prepared a Slurm job file (*pt_demo.slurm*) to run the PyTorch ResNet50 model. Submit the job using the command
- `$ cd pytorch/`
- `$ sbatch pt_demo.slurm`

Environment Setup for TensorFlow Models

Using the Intel AI Analytics Toolkit

```
# load all the necessary modules
module purge
module load intel/2023.07
module load intel/AIKit/2023.2.0

ENV_NAME=aikit-tf-gpu-clone

# If it doesn't exist, create the environment
if ! conda env list | grep -q "$ENV_NAME"; then

    conda create -n $ENV_NAME --clone aikit-tf-gpu
fi

# activate the conda environment
source activate $ENV_NAME
```

in tf_demo.slurm

Run Tensorflow ResNet50 Model

- We have prepared a Slurm job file (*tf_demo.slurm*) to run the Tensorflow ResNet50 model. Submit the job using the command

```
$ cd ..
```

```
$ cd tensorflow/
```

```
$ sbatch tf_demo.slurm
```

Lab III PyTorch on PVC



1. Import Intel Extension for PyTorch

Intel Extension for PyTorch is a Python package for extending PyTorch models to run on an Intel platform.

Add the following import statement to the beginning of your script:

```
import intel_extension_for_pytorch as ipex
```

2. Move the Model and Criterion to "xpu"

```
model = model.to("xpu")
```

```
criterion = criterion.to("xpu")
```


3. Apply the “ipex optimize” Function

Apply the ipex optimize function against the model and optimizer objects.

```
model, optimizer = ipex.optimize(model, optimizer=optimizer,  
dtype=torch.bfloat16)
```

4. Move the Data and Target to “xpu”

In the training loop,

```
data = data.to("xpu")
```

```
target = target.to("xpu")
```

5. Use Auto Mixed Precision (AMP)

Use automatic mixed-precision (AMP) with BFloat16 data type with the *torch.xpu.amp.autocast* context manager

with `torch.xpu.amp.autocast(enabled=True, dtype=torch.bfloat16):`

Hands-on Session

- Navigate to the pytorch exercises directory

```
$ cd $SCRATCH/aces_pvc_course/pytorch/exercises
```

- Open the exercise file (*cifar10_pvc_todo.py*) with your preferred editor (e.g. vim) or the file editor of the OnDemand portal.
- Complete the **#Todos** in the *cifar10_pvc_todo.py* file.
- Modify the Slurm job file (*pt_cifar10_pvc.slurm*) and submit your job.

```
$ sbatch pt_cifar10_pvc.slurm
```

Lab IV. TensorFlow on PVC



Install Intel Extension for Tensorflow

The Intel Extension for Tensorflow is based the on Tensorflow PluggableDevice interface to bring Intel XPU (GPU, CPU, etc) devices into Tensorflow.

To check the version, add import statement to the beginning of your script:

```
import intel_extension_for_tensorflow as itex
print(itex.__version__)
```

The default device will be Intel GPU after installing `intel-extension-for-tensorflow`

Source: Intel presentation at ACES Workshop

No Code Changes are Needed!



Credit: Bing Chat Enterprise

Hands-on Session

- Navigate to the pytorch exercises directory

```
$ cd $SCRATCH/aces_pvc_course/tensorflow/exercises
```

- Open the exercise file (*cifar10_pvc.py*) with your preferred editor (e.g. vim) or the file editor of the OnDemand portal.
- Read through the code to verify that there are no code changes
- Modify the Slurm job file (*tf_cifar10_pvc.slurm*) and submit your job.

```
$ sbatch tf_cifar10_pvc.slurm
```

```
##### TODO #####  
#uncomment below lines to run the cifar-10 exercise  
# cd $SCRATCH/aces_pvc_labs/tensorflow/exercises  
# python cifar10_pvc.py  
##### END TODO #####
```


PVC Monitoring Tools

- View the pvc nodes and number of GPUs

```
$ pestat -p pvc -G
```


- Monitor the system activity

```
$ watch -n 5 sysmon
```

- Intel XPU manager

```
$ watch -n 5 xpumcli stats -d <device index>
```

Start a VNC job



VNC Form

ACES OnDemand Portal Files Jobs Clusters Interactive Apps Affinity Groups Dashboard

Home / My Interactive Sessions / VNC

Interactive Apps

GUI

VNC

Nextsilicon VNC

Imaging

CryoSPARC

ImageJ

cisTEM

Servers

Jupyter Notebook

JupyterLab

RStudio

TensorBoard

Interactive Apps [Sandbox]

Servers

VNC

This app will launch a VNC job on ACES for remote visualization.

Node type

Intel GPU Max (PVC)

- select a non-CPU node type only if your software supports the Accelerator

Number of GPUs

1

- Current GPU Node Configuration
 - 13 x H100:2
 - 6 x PVC:4
 - 1 x H100:4
- Current GPU Node Availability
 - 5 x PVC:4
 - 1 x H100:1*
 - 1 x PVC:3*

Number of hours (max 48)

1

Number of cores (max 96)

3

Fields:

Node Type: Intel GPU Max (PVC)

Number of GPUs: 1

Number of hours: 1

Number of cores: 3

Total memory (GB): 5

Launch VNC

ACES OnDemand Portal Files Jobs Clusters Interactive Apps Affinity Groups Dashboard

Session was successfully deleted. X

Home / My Interactive Sessions

Interactive Apps

- GUI
- VNC
- Nextsilicon VNC
- Imaging
- CryoSPARC
- ImageJ
- cisTEM
- Servers
- Jupyter Notebook
- JupyterLab
- RStudio
- TensorBoard

VNC (28902) 1 node | 3 cores | Running

Host: >ac026 Delete

Created at: 2023-11-06 15:49:00 CST

Time Remaining: 56 minutes

Session ID: 83fc469d-4f99-4b4e-a8b1-39e6a343ba80

Compression 0 (low) to 9 (high) Image Quality 0 (low) to 9 (high)

Launch VNC View Only (Share-able Link)

My Interactive Sessions - ACE x TurboVNC: ac026:2 (u.zh1086 x +

portal-aces.hprc.tamu.edu/pun/sys/dashboard/noVNC-1.3.0/...

@ac026:/scratch/user/ /aces_pvc_course/pytorch/exercises

```
ses]$ pwd
/aces_pvc_course/pytorch/exercises
ses]$ ml purge
ses]$ ml intel/AIKit/2023.2.0
ses]$ ml intel/2023.07
ses]$ source activate aikit-pt-gpu-clone
(aikit-pt-gpu-clone) @ac026 exercises]$ python cifar10_pvc_solution.py > out.txt 2>&1 &
[1] 275895
(aikit-pt-gpu-clone) @ac026 exercises]$ watch sysmon
```

@ac026:/scratch/user/ /aces_p

Every 2.0s: sysmon

=====


GPU 0: Intel(R) Data Center GPU Max 1100 PCI Bus: 0000:21:00.0
Vendor: Intel(R) Corporation Driver Version: 1.3.26516 Subdevices: 0
EU Count: 448 Threads Per EU: 8 EU SIMD Width: 16 Total Memory(MB): 46679.2
Core Frequency(MHz): 200.0 of 1550.0 Core Temperature(C): unknown

=====

Running Processes: 3

PID,	Device Memory Used(MB),	Shared Memory Used(MB),	GPU Engines,	Executable
4793,	5.1,	0.0,	COMPUTE,	/usr/bin/xpumd
279895,	6888.2,	0.0,	COMPUTE;DMA,	python
280095,	1.8,	0.0,	UNKNOWN,	sysmon

Lab V. Running LAMMPS on PVC GPUs



LAMMPS on GPUs

- LAMMPS has a *modular* back-end for GPU acceleration

- GPU package

- OpenCL and CUDA GPUs are supported
 - divides work among the GPU and the CPUs

This one
today

- Kokkos package

- many more kinds of accelerators are supported
 - default strategy is “everything on the GPU”

LAMMPS on ACES

- Module `LAMMPS/3Aug2023-intel-2023.07`
- Requires toolchain module `intel/2023.07`
- uses GPU package
- Main executable is named `lmp_oneapi`
- Software located at `$HPRCROOTLAMMPS`

```
module load intel/2023.07 LAMMPS
echo $HPRCROOTLAMMPS
which lmp_oneapi
```


MPI on ACES

- Tip: don't mix `srun` and `mpirun` (or else performance very poor). Use `sbatch` with `mpirun` instead.
 - `sbatch filename`
 - `mpirun -np N`
- Suggested MPI thread configuration
 - `I_MPI_PIN_DOMAIN=auto:compact`
 - `I_MPI_FABRICS=shm`
 - `KMP_AFFINITY="granularity=core,scatter"`
 - `KMP_BLOCKTIME=1000`

Compute Aggregation Layer

- A tool for oneAPI Level Zero and OpenCL Applications.
- Helps to orchestrate efficient communication between MPI processes and GPU devices.
- `cal` provided with LAMMPS module.
- prepend an mpi launch with `calrun`

```
calrun mpirun -np N lmp_mpi
```

LAMMPS Testing

LAMMPS input files for common benchmarks

- `in.intel.airebo`
- `in.intel.dpd`
- `in.intel.eam`
- `in.intel.lc`
- `in.intel.lj`
- `in.intel.rhodo`
- `in.intel.snap`
- `in.intel.sw`
- `in.intel.tersoff`
- `in.intel.water`

```
cd $HPRCROOTLAMMPS/apps/TEST
```

LAMMPS Demo

- Exercise: try different input files.
- Exercise: try different resource combinations.
- Output files at `lammps_demo.<jobid>` and `demo-logs/<input>.log`

```
cd $SCRATCH/aces_pvc_course/lammps
cat lammps_demo.slurm
sbatch --reservation=training lammps_demo.slurm
```

Acknowledgement

This work was supported by the National Science Foundation (NSF)

- NSF award number 2112356 ACES - Accelerating Computing for Emerging Sciences,
- NSF award number 1925764 SWEETER - SouthWest Expertise in Expanding, Training, Education and Research,
- NSF award number 2019129 FASTER - Fostering Accelerated Scientific Transformations, Education, and Research,
- Dunni Aribuki from Intel,
- staff and students at Texas A&M High Performance Research Computing.



High Performance Research Computing

DIVISION OF RESEARCH

<https://hprc.tamu.edu>

HPRC Helpdesk:

help@hprc.tamu.edu

Phone: 979-845-0219

Help us help you. Please include details in your request for support, such as, Cluster (Faster, Grace, Terra, ViDaL), NetID (UserID), Job information (Job id(s), Location of your jobfile, input/output files, Application, Module(s) loaded, Error messages, etc), and Steps you have taken, so we can reproduce the problem.

Environment Setup for PyTorch models

Use Intel AI Analytics Toolkit

- Navigate to the pytorch directory

```
$ cd pytorch
```

- Clean up the environment and load the intel and intel/AIKit/2023.2.0 modules

```
$ module purge
```

```
$ module load intel/2023.07
```

```
$ module load intel/AIKit/2023.2.0
```

- List the pre-installed environments

```
$ conda env list
```

Environment Setup for PyTorch models

Use Intel AI Analytics Toolkit (Cont'd)

- Create a clone of the environment needed (aikit-pt-gpu for pytorch)

```
$ conda create -n aikit-pt-gpu-clone --clone aikit-pt-gpu
```

- Activate the environment

```
$ source activate aikit-pt-gpu-clone
```

- Install tqdm in the environment for the progress

```
$ conda install tqdm
```

- Deactivate the environment

```
$ conda deactivate
```